

Genome Sequencing & Assembly

Michael Schatz

Oct 23, 2014
Programming for Biology





Outline

1. Assembly theory
 1. Assembly by analogy
 2. De Bruijn and Overlap graph
 3. Coverage, read length, errors, and repeats
2. Whole Genome Alignment
 1. Aligning & visualizing with MUMmer
3. Genome assemblers
 1. ALLPATHS-LG: recommended for Illumina-only projects
 2. Celera Assembler: recommended for long read projects
4. Summary & Recommendations

Shredded Book Reconstruction

- Dickens accidentally shreds the first printing of A Tale of Two Cities
 - Text printed on 5 long spools

It was the best of times, it was the worst of times, it was the age of wisdom, it was the age of foolishness, ...

It was the best of times, it was the worst of times, it was the age of wisdom, it was the age of foolishness, ...

It was the best of times, it was the worst of times, it was the age of wisdom, it was the age of foolishness, ...

It was the best of times, it was the worst of times, it was the age of wisdom, it was the age of foolishness, ...

It was the best of times, it was the worst of times, it was the age of wisdom, it was the age of foolishness, ...

Greedy Reconstruction

The repeated sequence make the correct reconstruction ambiguous

- It was the best of times, it was the [worst/age]

Model the assembly problem as a graph problem

de Bruijn Graph Construction

- $D_k = (V, E)$
 - $V =$ All length- k subfragments ($k < l$)
 - $E =$ Directed edges between consecutive subfragments
 - Nodes overlap by $k-1$ words

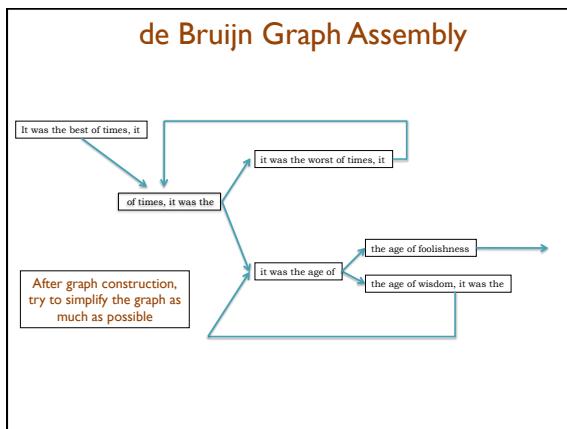
Original Fragment	Directed Edge
<code>[It was the best of]</code>	<code>[It was the best] → [was the best of]</code>

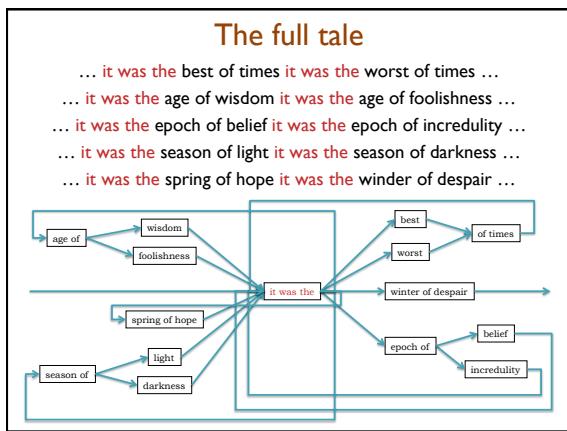
- Locally constructed graph reveals the global sequence structure
 - Overlaps between sequences implicitly computed

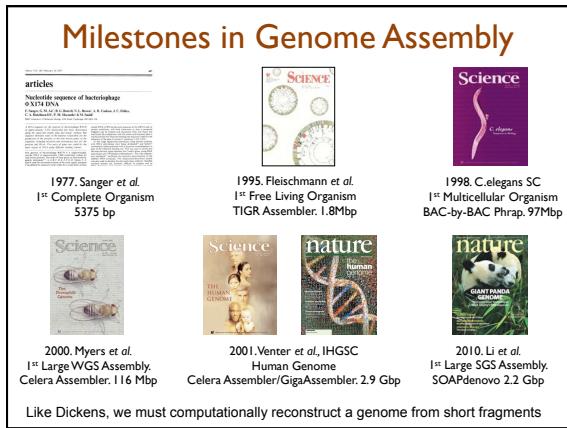
de Bruijn, 1946
Idury and Waterman, 1995
Pevzner, Tang, Waterman, 2001

de Bruijn Graph Assembly

After graph construction, try to simplify the graph as much as possible







Assembly Applications

- Novel genomes
- Metagenomes
- Sequencing assays
 - Structural variations
 - Transcript assembly
 - ...

The slide contains four logos: 'GENOME 10K' with a DNA double helix and elephant icon; 'i5k' with a butterfly icon; 'HMP' with a circular microbiome icon; and a sequencing assay visualization showing a stack of colored bars representing sequence reads.

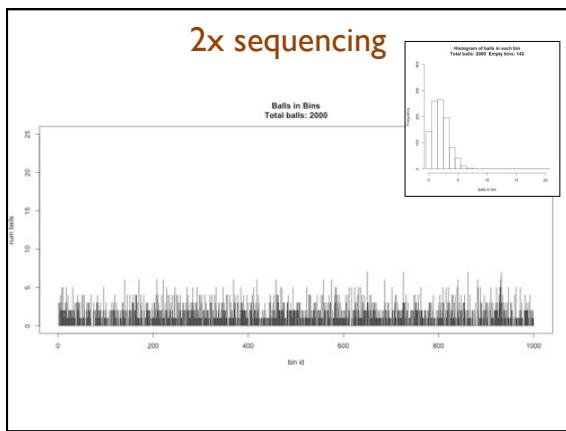
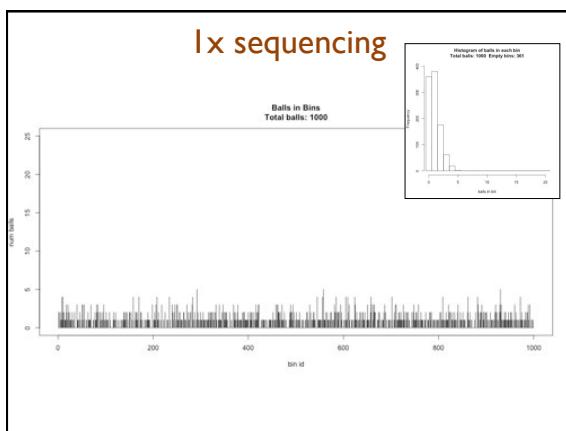
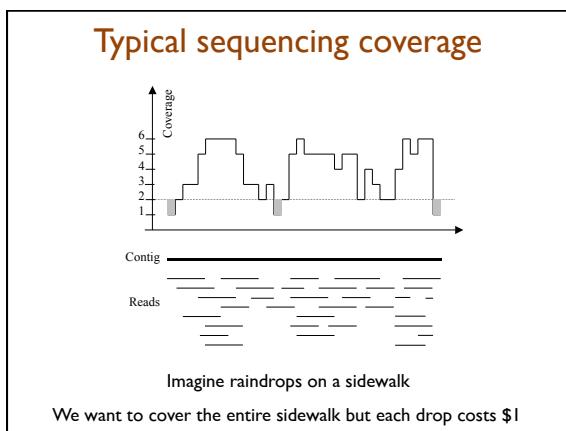
Ingredients for a good assembly

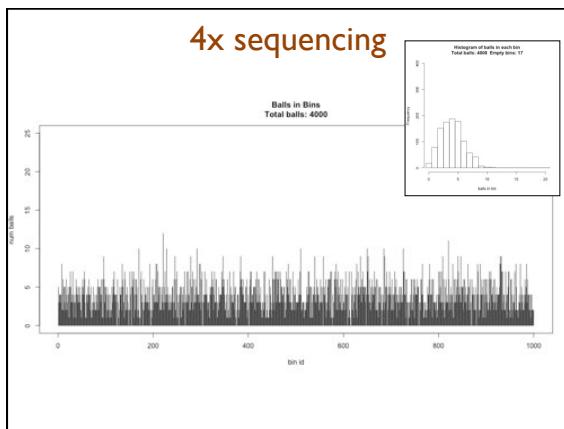
Coverage	Read Length	Quality
<p>High coverage is required</p> <ul style="list-style-type: none"> - Oversample the genome to ensure genome is sequenced with long overlaps between reads - Biased coverage will also fragment assembly 	<p>Reads & mates must be longer than the repeats</p> <ul style="list-style-type: none"> - Short repeats have false overlaps forming hairball assembly graphs - With long enough reads, assemble entire chromosomes into contigs 	<p>Errors obscure overlaps</p> <ul style="list-style-type: none"> - Reads are assembled by finding longer shared tail of reads - High error rates requires very short seeds, increasing complexity and forming assembly hairballs
Current challenges in de novo plant genome sequencing and assembly Schatz MC, Witkowski, McCombie, WR (2012) Genome Biology. 12:243		

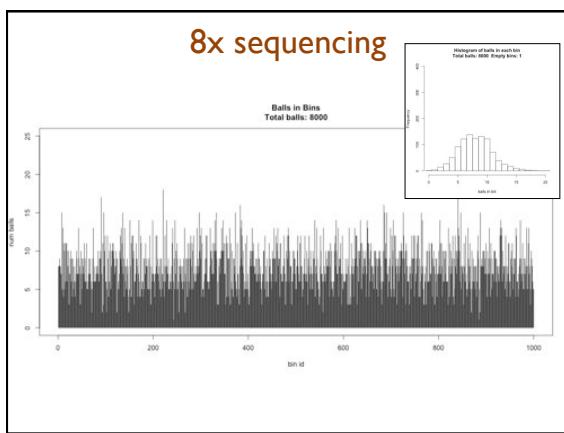
Illumina Sequencing by Synthesis

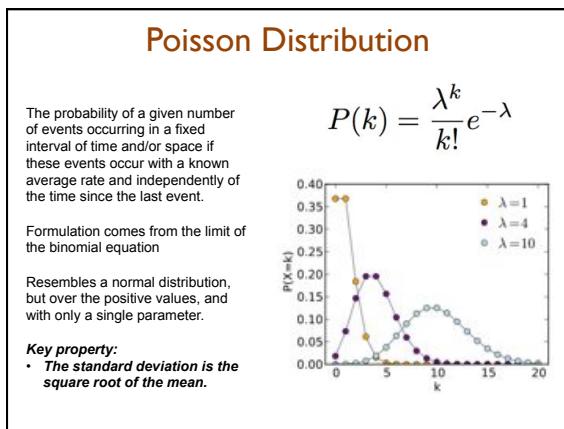
The diagram shows a sequential process: 1. Prepare (DNA fragments), 2. Attach (DNA fragments to a flowcell), 3. Amplify (using primers and nucleotides), 4. Image (laser illumination of the flowcell), and 5. Basecall (interpretation of image data).

Metzker (2010) Nature Reviews Genetics 11:31-46
<http://www.youtube.com/watch?v=l99aKKHcxC4>



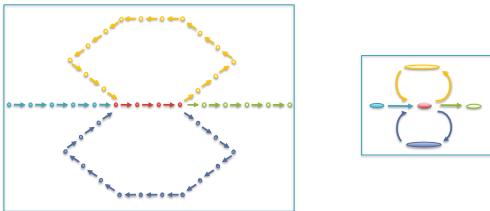






Unitigging / Unipathing

- After simplification and correction, compress graph down to its non-branching initial contigs
 - Aka "unitigs", "unipaths"
 - Unitigs end because of (1) lack of coverage, (2) errors, (3) heterozygosity, and (4) repeats

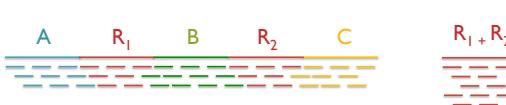


Repetitive regions

Repeat Type	Definition / Example	Prevalence
Low-complexity DNA / Microsatellites	$(b_1 b_2 \dots b_n)^N$ where $1 \leq k \leq 6$ CACACACACACACACACA	2%
SINEs (Short Interspersed Nuclear Elements)	Alu sequence (~280 bp) Mariner elements (~80 bp)	13%
LINEs (Long Interspersed Nuclear Elements)	~500 – 5,000 bp	21%
LTR (long terminal repeat) retrotransposons	Ty1-copia, Ty3-gypsy, Pao-BEL (~100 – 5,000 bp)	8%
Other DNA transposons		3%
Gene families & segmental duplications		4%

- Over 50% of mammalian genomes are repetitive
 - Large plant genomes tend to be even worse
 - Wheat: 16 Gbp; Pine: 24 Gbp

Repeats and Coverage Statistics

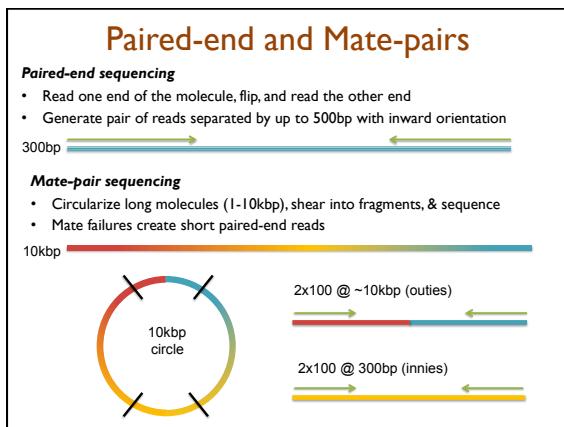


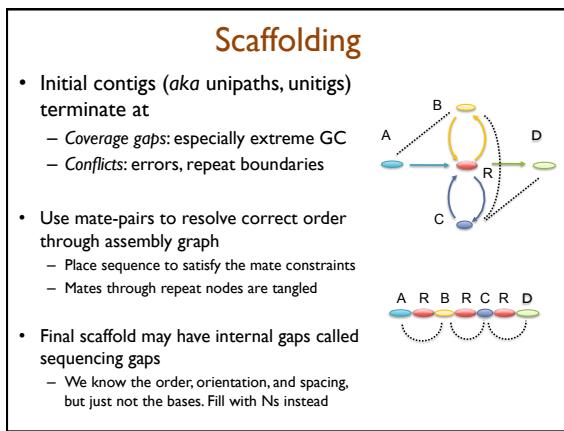
- If n reads are a uniform random sample of the genome of length G , we expect $k = n\Delta/G$ reads to start in a region of length Δ .
 - If we see many more reads than k (if the arrival rate is $> A$) , it is likely to be a collapsed repeat

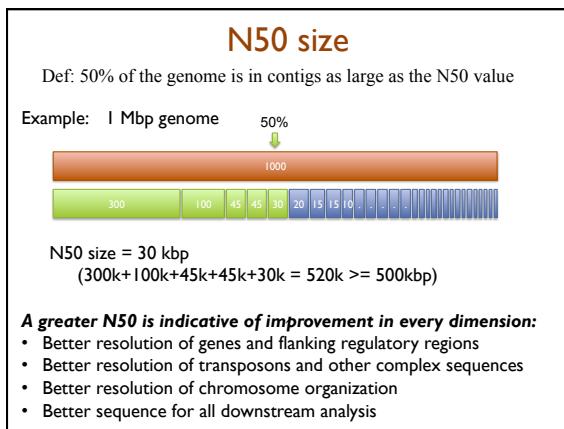
$$\Pr(X - \text{copy}) = \binom{n}{k} \left(\frac{X\Delta}{G}\right)^k \left(\frac{(G-X\Delta)}{G}\right)^{n-k}$$

$$A(\Delta, k) = \ln \left(\frac{\Pr(1 - \text{copy})}{\Pr(2 - \text{copy})} \right) = \ln \left(\frac{\left(\frac{(X\Delta/G)^k e^{-\Delta}}{k!}\right)}{\left(\frac{(2X\Delta/G)^k e^{-2\Delta}}{k!}\right)} \right) = \frac{n\Delta}{G} - k \ln 2$$

The fragment assembly string graph
Myers, EW (2005) Bioinformatics. 21(suppl 2):ii79-85.









Outline

1. Assembly theory
 1. Assembly by analogy
 2. De Bruijn and Overlap graph
 3. Coverage, read length, errors, and repeats
2. Whole Genome Alignment
 1. Aligning & visualizing with MUMmer
3. Genome assemblers
 1. ALLPATHS-LG: recommended for Illumina-only projects
 2. Celera Assembler: recommended for long read projects
4. Summary and Recommendations

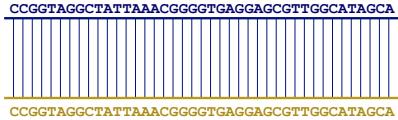


Whole Genome Alignment with MUMmer

Slides Courtesy of Adam M. Phillippy
University of Maryland

Goal of WGA

- For two genomes, A and B, find a mapping from each position in A to its corresponding position in B



```
CCGGTAGGCTATTAAACGGGGTGAGGGAGCGTTGGCATAGCA
          |-----|-----|-----|-----|-----|-----|
CCGGTAGGCTATTAAACGGGGTGAGGGAGCGTTGGCATAGCA
```

Not so fast...

- Genome A may have insertions, deletions, translocations, inversions, duplications or SNPs with respect to B (sometimes all of the above)



WGA visualization

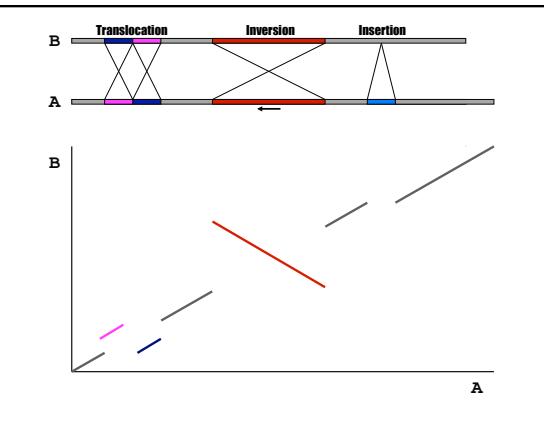
- How can we visualize *whole genome* alignments?

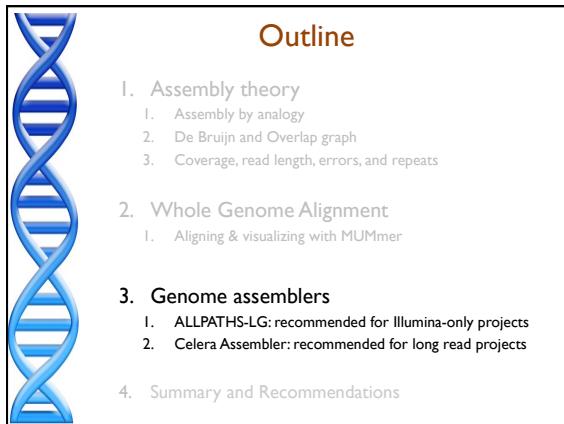
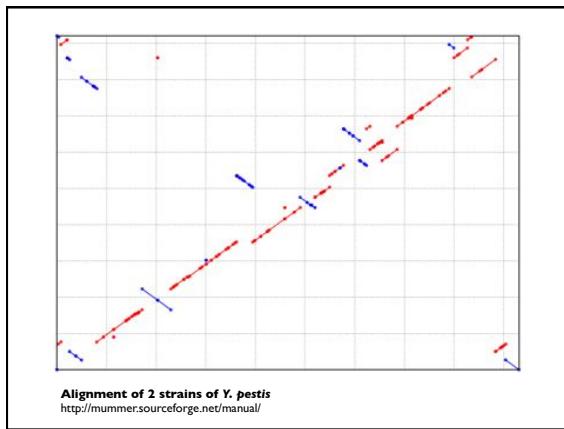
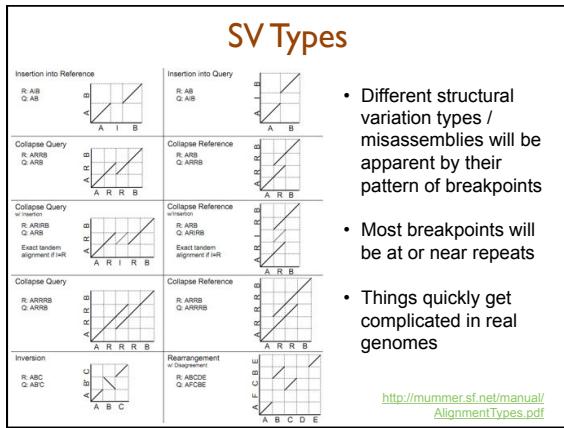
- With an alignment dot plot

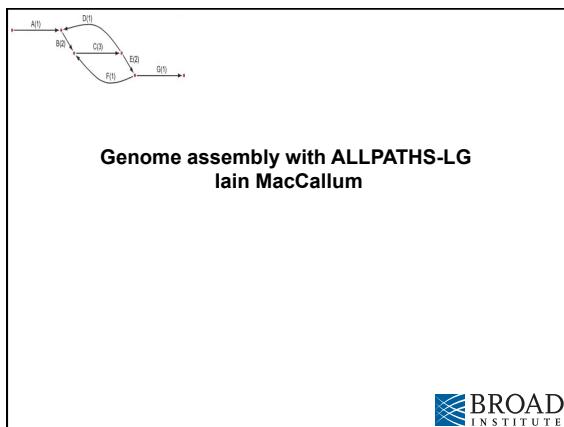
T				
G				
C				
A				

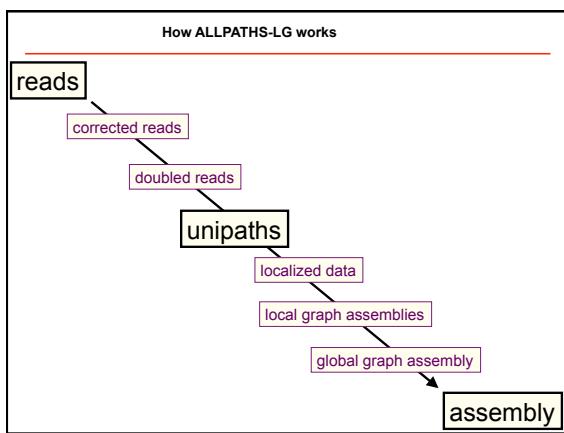
A C C T

- N × M matrix
 - Let i = position in genome A
 - Let j = position in genome B
 - Fill cell (i,j) if A_i shows similarity to B_j









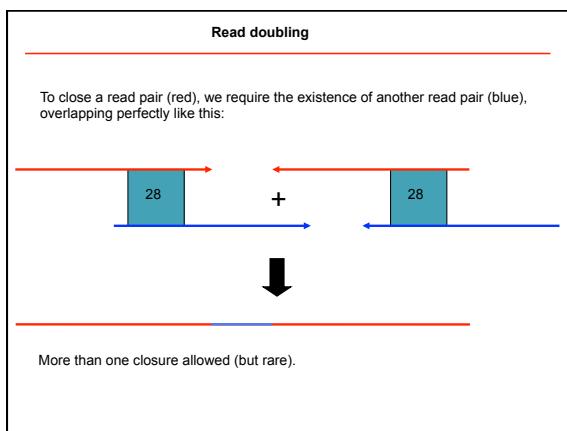
ALLPATHS-LG sequencing model

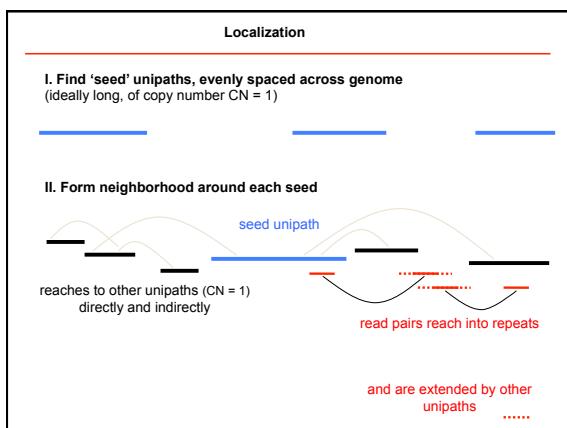
Libraries (insert types)	Fragment size (bp)	Read length (bases)	Sequence coverage (x)	Required
Fragment	180*	≥ 100	45	yes
Short jump	3,000	≥ 100 preferable	45	yes
Long jump	6,000	≥ 100 preferable	5	no**
Fosmid jump	40,000	≥ 26	1	no**

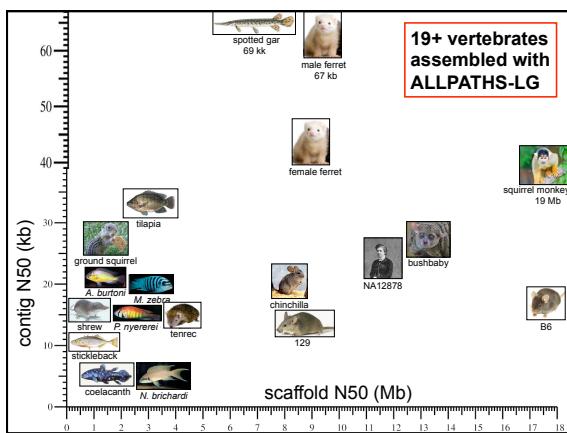
*See next slide.
**For best results. Normally not used for small genomes.
However essential to assemble long repeats or duplications.

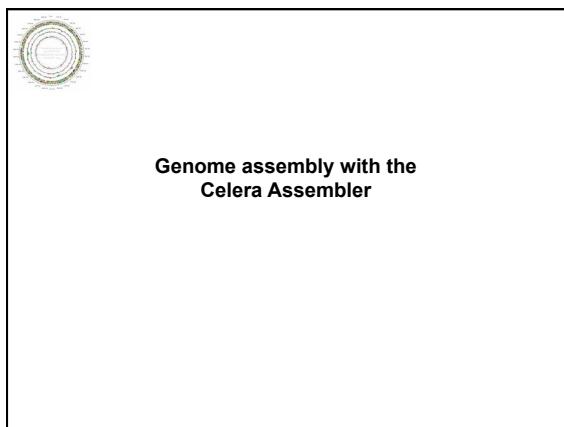
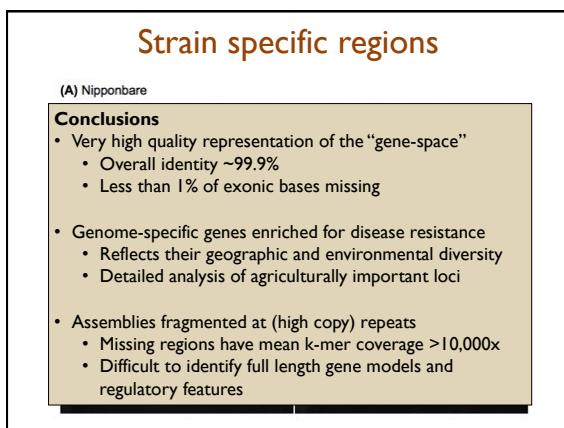
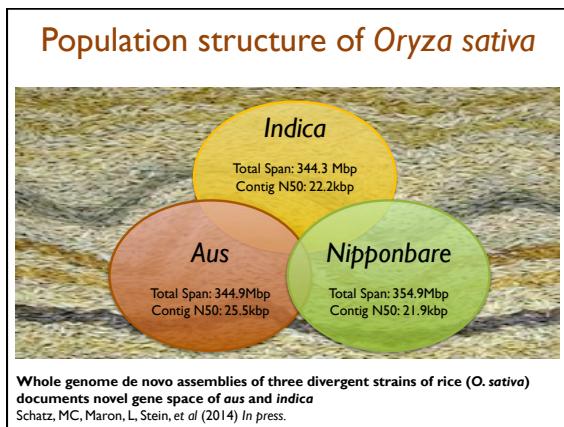
Cutting coverage in half still works, with some reduction in quality of results.

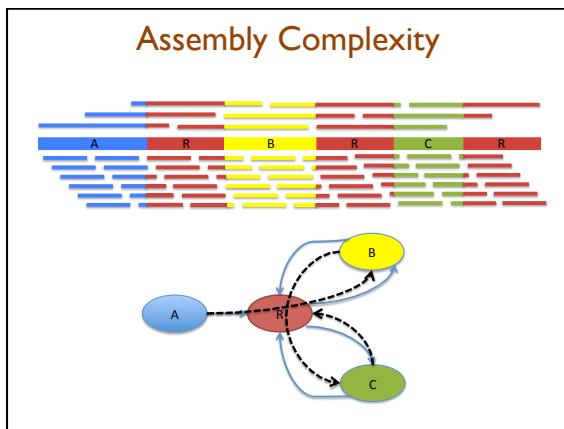
All: protocols are either available, or in progress.

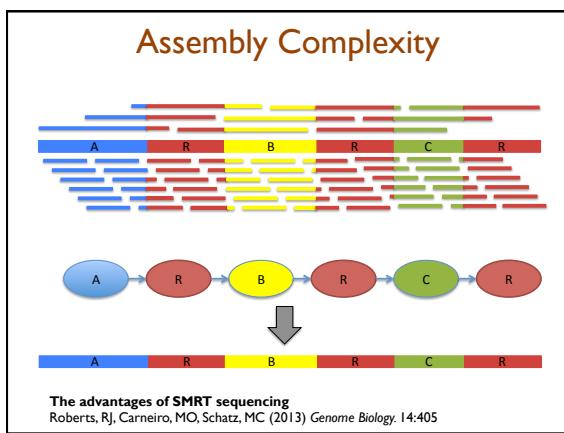


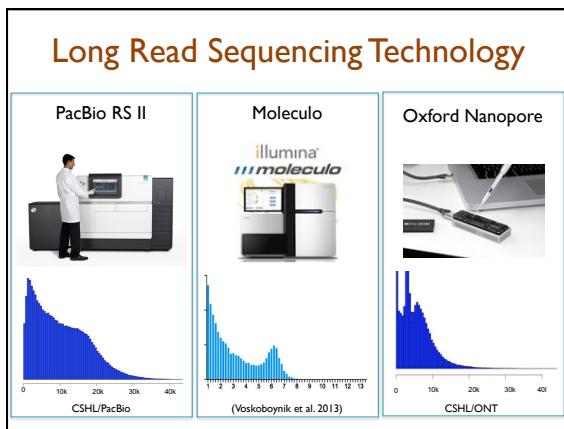


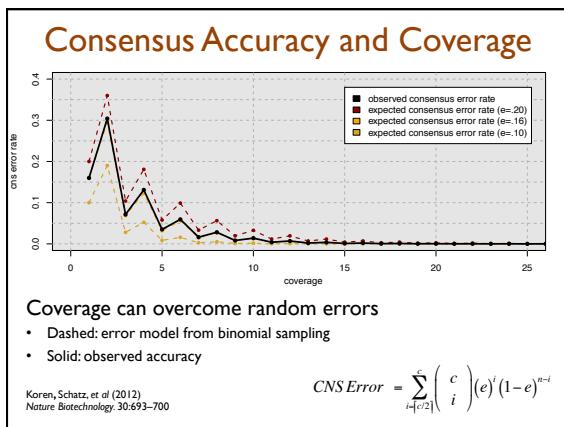
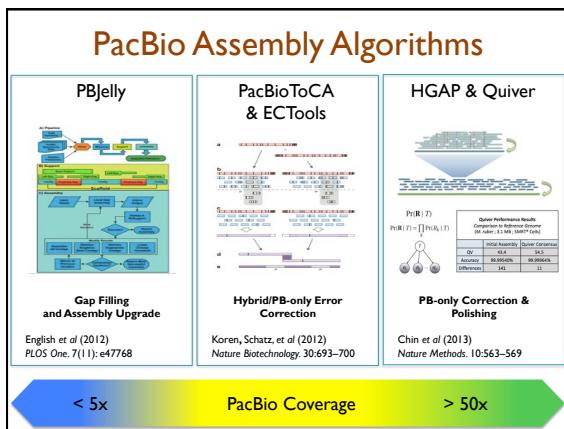
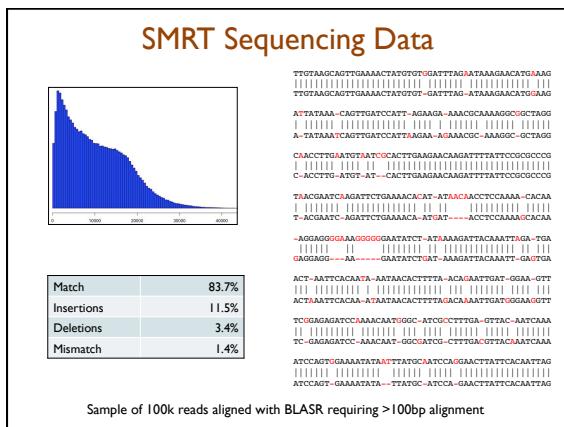


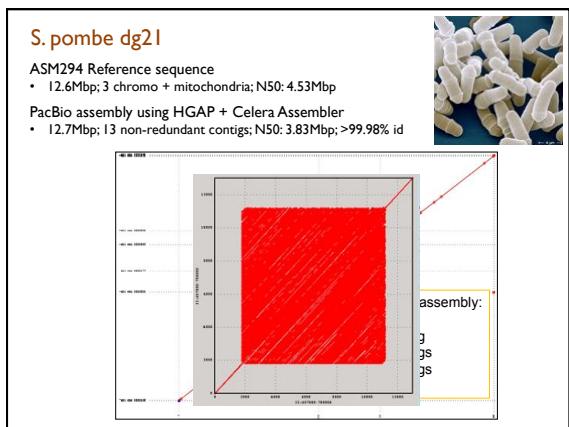
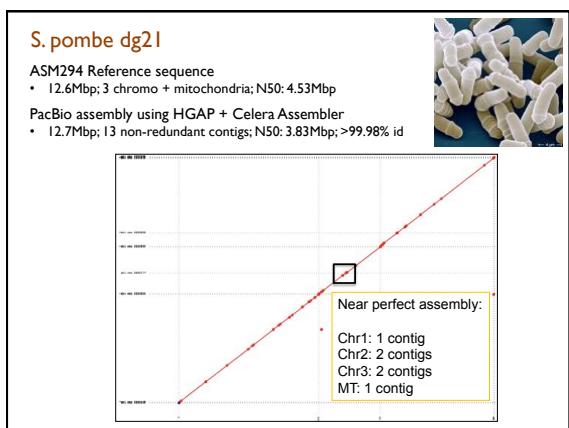
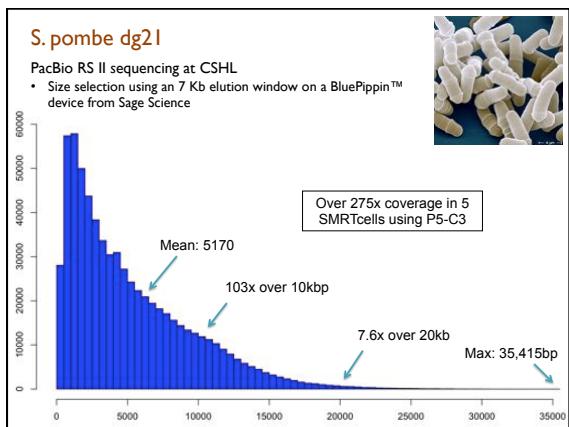












A. thaliana Ler-0

<http://blog.pacificbiosciences.com/2013/08/new-data-release-arabidopsis-assembly.html>



A. thaliana Ler-0 sequenced at PacBio

- Sequenced using the previous P4 enzyme and C2 chemistry
- Size selection using an 8 Kb to 50 Kb elution window on a BluePippin™ device from Sage Science
- Total coverage >119x

Genome size:	124.6 Mbp	Sum of Contig Lengths:	149.5Mb
Chromosome N50:	23.0 Mbp	N50 Contig Length:	8.4 Mb
Corrected coverage:	20x over 10kb	Number of Contigs:	1788

High quality assembly of chromosome arms
 Assembly Performance: 8.4Mbp/23Mbp = 36%
 MiSeq assembly: 63kbp/23Mbp = .2%

Human CHM1

<http://blog.pacificbiosciences.com/2014/02/data-release-54x-long-read-coverage-for.html>



CHM1 hert sequenced at PacBio

- Sequenced using the P5 enzyme and C3 chemistry
- Size selection using an 20kb elution window on a BluePippin™ device from Sage Science
- Total coverage: 54x

Genome size:	3.0 Gb	Sum of Contig Lengths:	3.2 Gb
Chromosome N50:	90.5 Mbp	N50 Contig Length:	4.38 Mbp
Average read length:	7,680 bp	Max Contig:	44 Mbp

High quality draft assembly
 Assembly Performance: 4.38Mbp/90.5Mbp = 4.5%
 Sanger HuRef assembly: 107kbp / 90.5Mbp = .1%

Current Collaborations



Indica & Aus Rice
McCombie/Ware/McCouch



Pineapple
UIUC



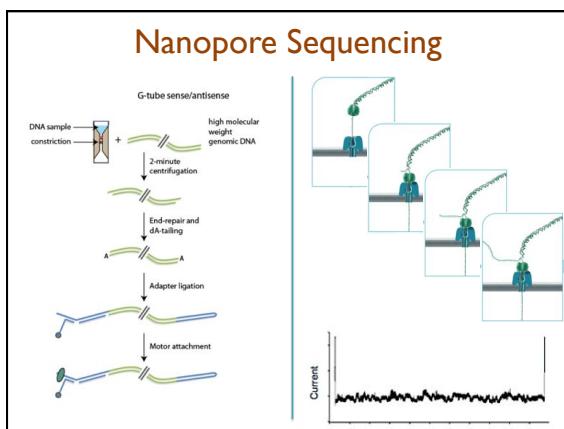
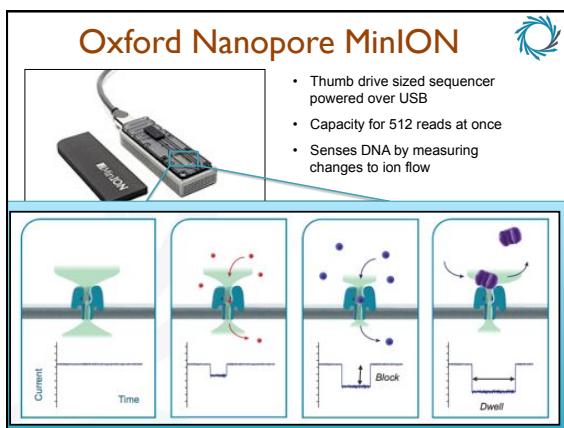
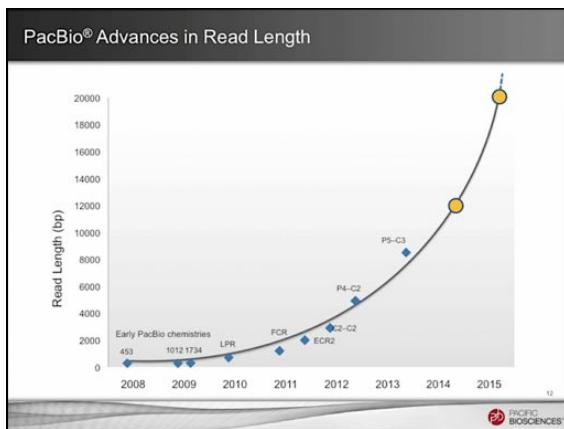
M. ligano
Hannon

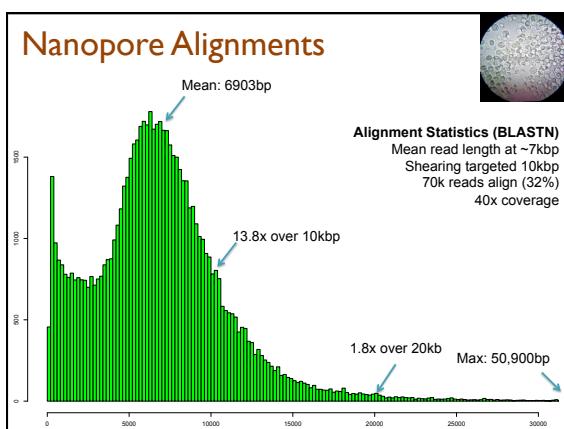
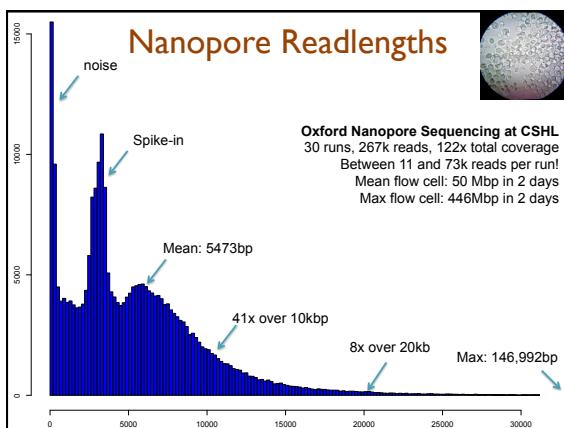
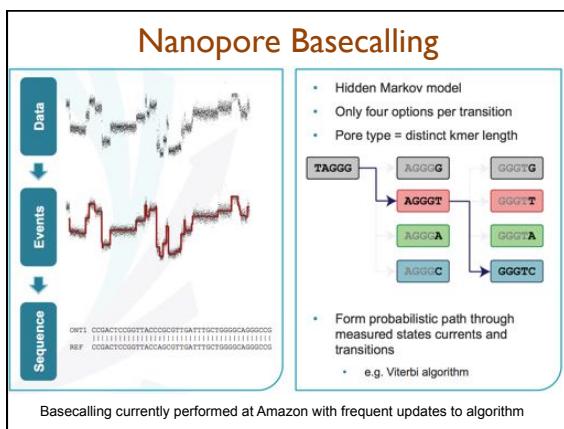


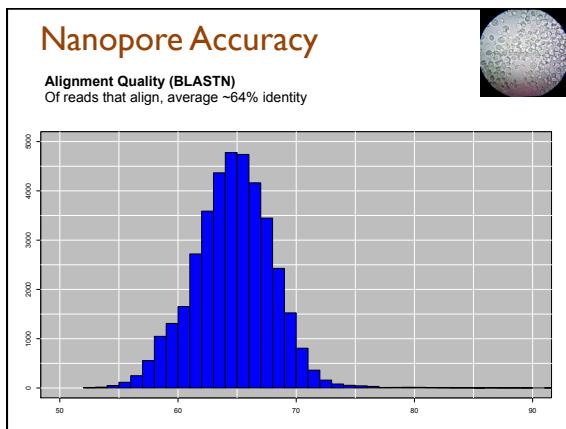
Asian Sea Bass
Temasek Life Sciences Laboratory

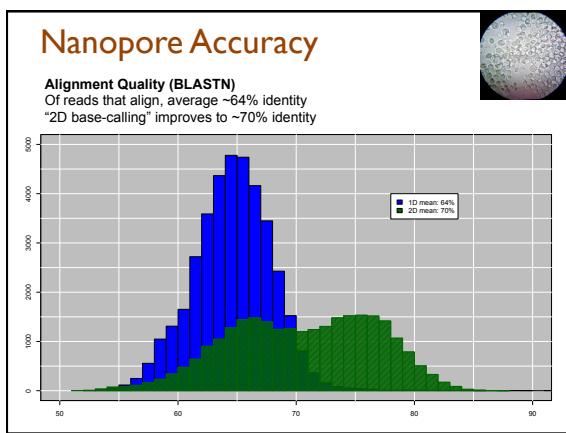


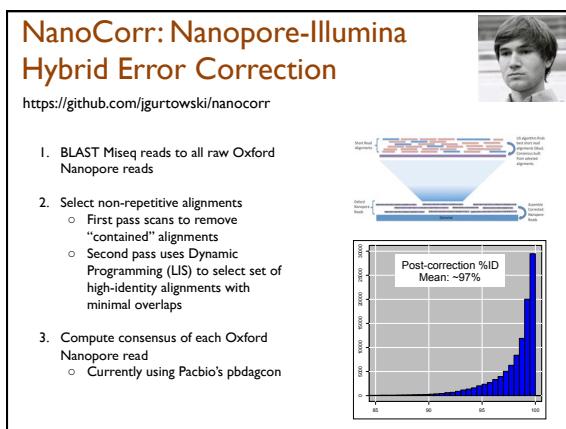
P. hominis
NYU

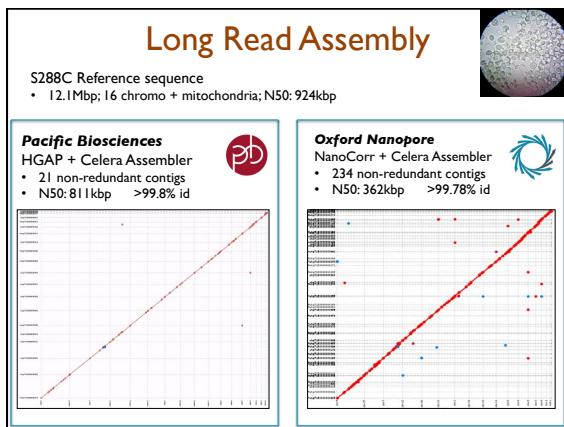














What should we expect from an assembly?

Analysis of dozens of genomes from across the tree of life with real and simulated data

Summary & Recommendations

- < 100 Mbp: HGAP/PacBio2CA @ 100x PB C3-P5
expect near perfect chromosome arms
- < 1GB: HGAP/PacBio2CA @ 100x PB C3-P5
high quality assembly: contig N50 over 1Mbp
- > 1GB: hybrid/gap filling
expect contig N50 to be 100kbp – 1Mbp
- > 5GB: Email mschatz@cshl.edu

Error correction and assembly complexity of single molecule sequencing reads.
 Lee, H^a, Gurtowski, J^a, Yoo, S, Marcus, S, McCombie, VR, Schatz, MC
<http://www.biorxiv.org/content/early/2014/06/18/006395>

Assembly Summary



Assembly quality depends on

1. **Coverage:** low coverage is mathematically hopeless
2. **Repeat composition:** high repeat content is challenging
3. **Read length:** longer reads help resolve repeats
4. **Error rate:** errors reduce coverage, obscure true overlaps

- Assembly is a hierarchical, starting from individual reads, build high confidence contigs/unitigs, incorporate the mates to build scaffolds
 - Extensive error correction is the key to getting the best assembly possible from a given data set
- Watch out for collapsed repeats & other misassemblies
 - Globally/Locally reassemble data from scratch with better parameters & stitch the 2 assemblies together

Acknowledgements

Schatz Lab Rahul Amin Tyler Gavin James Gurtowski Han Fang Hayan Lee Maria Nattestad Aspyn Palatnick Srividya Ramakrishnan Eric Biggers Ke Jiang Shoshana Marcus Giuseppe Narzisi Rachel Sherman Greg Vrtute Alejandro Wences	<table border="0" style="width: 100%; border-collapse: collapse;"> <tr> <td style="width: 33%; text-align: center;"> CSHL  </td> <td style="width: 33%; text-align: center;"> NSF  </td> <td style="width: 33%; text-align: center;"> National Human Genome Research Institute  </td> </tr> <tr> <td style="text-align: center;"> U.S. DEPARTMENT OF ENERGY  </td> <td colspan="2" style="text-align: center;"> SFARI SIMONS FOUNDATION AUTISM RESEARCH INITIATIVE </td> </tr> </table>	CSHL 	NSF 	National Human Genome Research Institute 	U.S. DEPARTMENT OF ENERGY 	SFARI SIMONS FOUNDATION AUTISM RESEARCH INITIATIVE	
CSHL 	NSF 	National Human Genome Research Institute 					
U.S. DEPARTMENT OF ENERGY 	SFARI SIMONS FOUNDATION AUTISM RESEARCH INITIATIVE						

Biological Data Sciences
 Anne Carpenter, Michael Schatz, Matt Wood
 Nov 5 - 8, 2014



Thank you

<http://schatzlab.cshl.edu>
 @mike_schatz
